

```
In [1]: import pandas as pd
```

```
In [2]: # Load the dataset
file_path = r'C:\Users\zahus\Desktop\DATA science\Module 11-Dissertaion\Dataset\6-Sentiment Analysis Dataset (Customer Feedback)'
df = pd.read_csv(file_path, delimiter=';')
```

```
In [3]: # Display basic information and the first few rows
df_info = df.info()
df_head = df.head()

df_info, df_head
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 23486 entries, 0 to 23485
Data columns (total 11 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Unnamed: 0                            23486 non-null  int64
1   Clothing.ID                           23486 non-null  int64
2   Age                                    23486 non-null  int64
3   Title                                 19676 non-null  object
4   Review.Text                           22641 non-null  object
5   Rating                                23486 non-null  int64
6   Recommended.IND                       23486 non-null  int64
7   Positive.Feedback.Count               23486 non-null  int64
8   Division.Name                         23472 non-null  object
9   Department.Name                       23472 non-null  object
10  Class.Name                            23472 non-null  object
dtypes: int64(6), object(5)
memory usage: 2.0+ MB
```

```
Out[3]: (None,
         Unnamed: 0  Clothing.ID  Age  Title \
0          1          767  33      NaN
1          2         1080  34      NaN
2          3         1077  60  Some major design flaws
3          4         1049  50      My favorite buy!
4          5          847  47      Flattering shirt

         Review.Text  Rating  Recommended.IND \
0  Absolutely wonderful - silky and sexy and comf...      4      1
1  Love this dress! it's sooo pretty. i happene...      5      1
2  I had such high hopes for this dress and reall...      3      0
3  I love, love, love this jumpsuit. it's fun, fl...      5      1
4  This shirt is very flattering to all due to th...      5      1

         Positive.Feedback.Count  Division.Name  Department.Name  Class.Name
0          0          Initmates      Intimate  Intimates
1          4          General      Dresses  Dresses
2          0          General      Dresses  Dresses
3          0  General Petite      Bottoms    Pants
4          6          General      Tops    Blouses )
```

Step 1- Clean and Preprocesse the Dataset

*Removed missing values. *Created sentiment labels from ratings. *Cleaned text by removing punctuation, numbers, and converting to lowercase

```
In [4]: # Keep only the necessary columns for sentiment analysis
df_clean = df[['Review.Text', 'Rating']].dropna()
```

```
In [5]: # Define sentiment from rating
def sentiment_from_rating(rating):
    if rating <= 2:
        return 'negative'
    elif rating == 3:
        return 'neutral'
    else:
        return 'positive'
```


```
In [6]: df_clean['Sentiment'] = df_clean['Rating'].apply(sentiment_from_rating)
```



Step 2 - Exploratory Data Analysis (EDA)

```
In [7]: import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [8]: # Plot sentiment distribution
plt.figure(figsize=(6, 4))
sns.countplot(data=df_clean, x='Sentiment', order=['negative', 'neutral', 'positive'])
plt.title('Sentiment Distribution Based on Ratings')
plt.xlabel('Sentiment')
plt.ylabel('Number of Reviews')
plt.tight_layout()
plt.show()
```



```
In [9]: import re
```

```
In [10]: # Clean review text for analysis
```

```
def clean_text(text):  
    text = str(text).lower()  
    text = re.sub(r"^[a-zA-Z\s]", "", text)  
    text = re.sub(r"\s+", " ", text).strip()  
    return text
```


```
df_clean['Clean_Review'] = df_clean['Review.Text'].apply(clean_text)
```

```
In [11]: # Plot distribution of review lengths
```

```
df_clean['Review_Length'] = df_clean['Clean_Review'].apply(lambda x: len(x.split()))
```

```
plt.figure(figsize=(8, 4))  
sns.histplot(data=df_clean, x='Review_Length', bins=40, kde=True)
```

```
plt.title('Distribution of Review Lengths')  
plt.xlabel('Number of Words')  
plt.ylabel('Frequency')  
plt.tight_layout()  
plt.show()
```



✓ Step 3 - Sentiment Encoding & Splitting

Sentiments are encoded as:

0 = Negative

1 = Neutral

2 = Positive

Data split into:

Training set: 18,112 reviews

Test set: 4,529 reviews

```
In [12]: from sklearn.preprocessing import LabelEncoder  
from sklearn.model_selection import train_test_split
```

```
In [13]: # Step 3: Encode sentiment labels  
label_encoder = LabelEncoder()  
df_clean['Sentiment_Label'] = label_encoder.fit_transform(df_clean['Sentiment'])
```

```
In [14]: # Split data into train and test sets  
X = df_clean['Clean_Review']  
y = df_clean['Sentiment_Label']  
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42, stratify=y)
```

```
In [15]: # Show label encoding and shapes  
label_mapping = dict(zip(label_encoder.classes_, label_encoder.transform(label_encoder.classes_)))  
X_train.shape, X_test.shape, y_train.shape, y_test.shape, label_mapping
```

```
Out[15]: ((18112,),  
(4529,),  
(18112,),  
(4529,),  
{'negative': 0, 'neutral': 1, 'positive': 2})
```



4: Train multiple NLP models (Logistic Regression, Naive Bayes, XGBoost).

Logistic Regression

```
In [16]: from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import classification_report
```

```
In [17]: # Vectorise text using TF-IDF
tfidf = TfidfVectorizer(max_features=10000)
X_train_tfidf = tfidf.fit_transform(X_train)
X_test_tfidf = tfidf.transform(X_test)
```

```
In [18]: # Train Logistic Regression
logreg = LogisticRegression(max_iter=1000)
logreg.fit(X_train_tfidf, y_train)
y_pred_logreg = logreg.predict(X_test_tfidf)
```

```
In [19]: # Classification report
logreg_report = classification_report(y_test, y_pred_logreg, target_names=label_mapping.keys(), output_dict=True)
logreg_report
```

```
Out[19]: {'negative': {'precision': 0.6056338028169014,
  'recall': 0.45358649789029537,
  'f1-score': 0.5186972255729795,
  'support': 474},
  'neutral': {'precision': 0.4477124183006536,
  'recall': 0.2424778761061947,
  'f1-score': 0.31458094144661314,
  'support': 565},
  'positive': {'precision': 0.8733195449844882,
  'recall': 0.9679083094555874,
  'f1-score': 0.9181842892090242,
  'support': 3490},
  'accuracy': 0.8235813645396335,
  'macro avg': {'precision': 0.6422219220340144,
  'recall': 0.5546575611506924,
  'f1-score': 0.5838208187428723,
  'support': 4529},
  'weighted avg': {'precision': 0.7922086886445008,
  'recall': 0.8235813645396335,
  'f1-score': 0.8010739426315794,
  'support': 4529}}
```

Naive Bayes model

```
In [20]: from sklearn.naive_bayes import MultinomialNB
```

```
In [21]: # Train Naive Bayes
nb = MultinomialNB()
nb.fit(X_train_tfidf, y_train)
y_pred_nb = nb.predict(X_test_tfidf)
```

```
In [22]: # Classification report
nb_report = classification_report(y_test, y_pred_nb, target_names=label_mapping.keys(), output_dict=True)
nb_report
```

```
Out[22]: {'negative': {'precision': 1.0,
                        'recall': 0.008438818565400843,
                        'f1-score': 0.01673640167364017,
                        'support': 474},
          'neutral': {'precision': 0.0, 'recall': 0.0, 'f1-score': 0.0, 'support': 565},
          'positive': {'precision': 0.7714412024756853,
                       'recall': 1.0,
                       'f1-score': 0.8709757923633641,
                       'support': 3490},
          'accuracy': 0.7714727312872599,
          'macro avg': {'precision': 0.5904804008252285,
                        'recall': 0.3361462728551336,
                        'f1-score': 0.29590406467900143,
                        'support': 4529},
          'weighted avg': {'precision': 0.6991233819033211,
                           'recall': 0.7714727312872599,
                           'f1-score': 0.6729164428662942,
                           'support': 4529}}
```

XGBoost Model

```
In [23]: import xgboost as xgb
```



```
In [24]: # Train XGBoost Classifier
xgb_clf = xgb.XGBClassifier(use_label_encoder=False, eval_metric='mlogloss')
xgb_clf.fit(X_train_tfidf, y_train)
y_pred_xgb = xgb_clf.predict(X_test_tfidf)

In [25]: # Classification report
xgb_report = classification_report(y_test, y_pred_xgb, target_names=label_mapping.keys(), output_dict=True)
xgb_report

Out[25]: {'negative': {'precision': 0.5901060070671378,
  'recall': 0.35232067510548526,
  'f1-score': 0.441215323645971,
  'support': 474},
  'neutral': {'precision': 0.41544117647058826,
  'recall': 0.2,
  'f1-score': 0.2700119474313023,
  'support': 565},
  'positive': {'precision': 0.8558127830900856,
  'recall': 0.9744985673352435,
  'f1-score': 0.9113076098606645,
  'support': 3490},
  'accuracy': 0.8127621991609627,
  'macro avg': {'precision': 0.6204533222092706,
  'recall': 0.508939747480243,
  'f1-score': 0.540844960312646,
  'support': 4529},
  'weighted avg': {'precision': 0.7730671505939731,
  'recall': 0.8127621991609627,
  'f1-score': 0.7821056242262299,
  'support': 4529}}
```



Visual comparison of Logistic Regression and Naive Bayes models

from sklearn.linear_model import LogisticRegression

```
from sklearn.naive_bayes import MultinomialNB from sklearn.metrics import classification_report
```

```
In [26]: # Train and evaluate models
logreg = LogisticRegression(max_iter=1000)
logreg.fit(X_train_tfidf, y_train)
y_pred_logreg = logreg.predict(X_test_tfidf)
logreg_report = classification_report(y_test, y_pred_logreg, output_dict=True, zero_division=0)


nb = MultinomialNB()
nb.fit(X_train_tfidf, y_train)
y_pred_nb = nb.predict(X_test_tfidf)
nb_report = classification_report(y_test, y_pred_nb, output_dict=True, zero_division=0)

In [27]: xgb_report = classification_report(y_test, y_pred_xgb, target_names=["negative", "neutral", "positive"], output_dict=True)
logreg_report = classification_report(y_test, y_pred_logreg, target_names=["negative", "neutral", "positive"], output_dict=True)
nb_report = classification_report(y_test, y_pred_nb, target_names=["negative", "neutral", "positive"], output_dict=True)

In [28]: # Create DataFrame
results = pd.DataFrame([
    {
        "Model": "Logistic Regression",
        "Accuracy": logreg_report["accuracy"],
        "F1_Negative": logreg_report["negative"]["f1-score"],
        "F1_Neutral": logreg_report["neutral"]["f1-score"],
        "F1_Positive": logreg_report["positive"]["f1-score"],
        "Macro_F1": logreg_report["macro avg"]["f1-score"]
    },
    {
        "Model": "Naive Bayes",
        "Accuracy": nb_report["accuracy"],
        "F1_Negative": nb_report["negative"]["f1-score"],
        "F1_Neutral": nb_report["neutral"]["f1-score"],
        "F1_Positive": nb_report["positive"]["f1-score"],
        "Macro_F1": nb_report["macro avg"]["f1-score"]
    },
    {
        "Model": "XGBoost",
        "Accuracy": xgb_report["accuracy"],
        "F1_Negative": xgb_report["negative"]["f1-score"],
```

```
    "F1_Neutral": xgb_report["neutral"]["f1-score"],  
    "F1_Positive": xgb_report["positive"]["f1-score"],  
    "Macro_F1": xgb_report["macro avg"]["f1-score"]  
}  
])
```

```
In [29]: # Melt and plot  
results_melted = results.melt(id_vars='Model', var_name='Metric', value_name='Score')  
plt.figure(figsize=(10, 6))  
sns.barplot(data=results_melted, x='Model', y='Score', hue='Metric')  
plt.title('Logistic Regression vs Naive Bayes vs XGBoost')  
plt.ylabel('Score')  
plt.ylim(0, 1)  
plt.xticks(rotation=15)  
plt.legend(bbox_to_anchor=(1.05, 1), loc='upper left')  
plt.tight_layout()  
plt.show()
```



Forecasting for 5 years

```
In [30]: from prophet import Prophet
import matplotlib.pyplot as plt
```

```
In [31]: # Create Sentiment column
def sentiment_from_rating(r):
    if r <= 2:
        return 'negative'
    elif r == 3:
        return 'neutral'
    else:
        return 'positive'
df['Sentiment'] = df['Rating'].apply(sentiment_from_rating)
```

```
In [32]: # Simulate a date column for each review (e.g., starting from Jan 2015)
start_date = pd.to_datetime("2015-01-01")
df['Date'] = pd.date_range(start=start_date, periods=len(df), freq='D')
```

```
In [33]: # Group by month and sentiment
df['Month'] = df['Date'].dt.to_period('M').dt.to_timestamp()
monthly_sentiment = df.groupby(['Month', 'Sentiment']).size().unstack().fillna(0)
```

```
In [34]: # Forecasting for each sentiment
forecast_years = 5
forecast_horizon = forecast_years * 12 # 5 years in months

for sentiment in ['positive', 'neutral', 'negative']:
    sentiment_df = monthly_sentiment[[sentiment]].reset_index()
    sentiment_df.columns = ['ds', 'y'] # Prophet requires 'ds' and 'y' columns


    model = Prophet()
    model.fit(sentiment_df)

    future = model.make_future_dataframe(periods=forecast_horizon, freq='M')
    forecast = model.predict(future)
```

```
19:54:58 - cmdstanpy - INFO - Chain [1] start processing
19:54:58 - cmdstanpy - INFO - Chain [1] done processing
19:54:58 - cmdstanpy - INFO - Chain [1] start processing
19:54:58 - cmdstanpy - INFO - Chain [1] done processing
19:54:59 - cmdstanpy - INFO - Chain [1] start processing
19:54:59 - cmdstanpy - INFO - Chain [1] done processing
```

In [36]:

```
# Plot forecast
fig = model.plot(forecast)
plt.title(f'5-Year Forecast for {sentiment.capitalize()} Sentiment')
plt.xlabel("Date")
plt.ylabel("Number of Reviews")
plt.tight_layout()
plt.show()
```



In []: